

Management and Operations of Networks, Services, and Systems

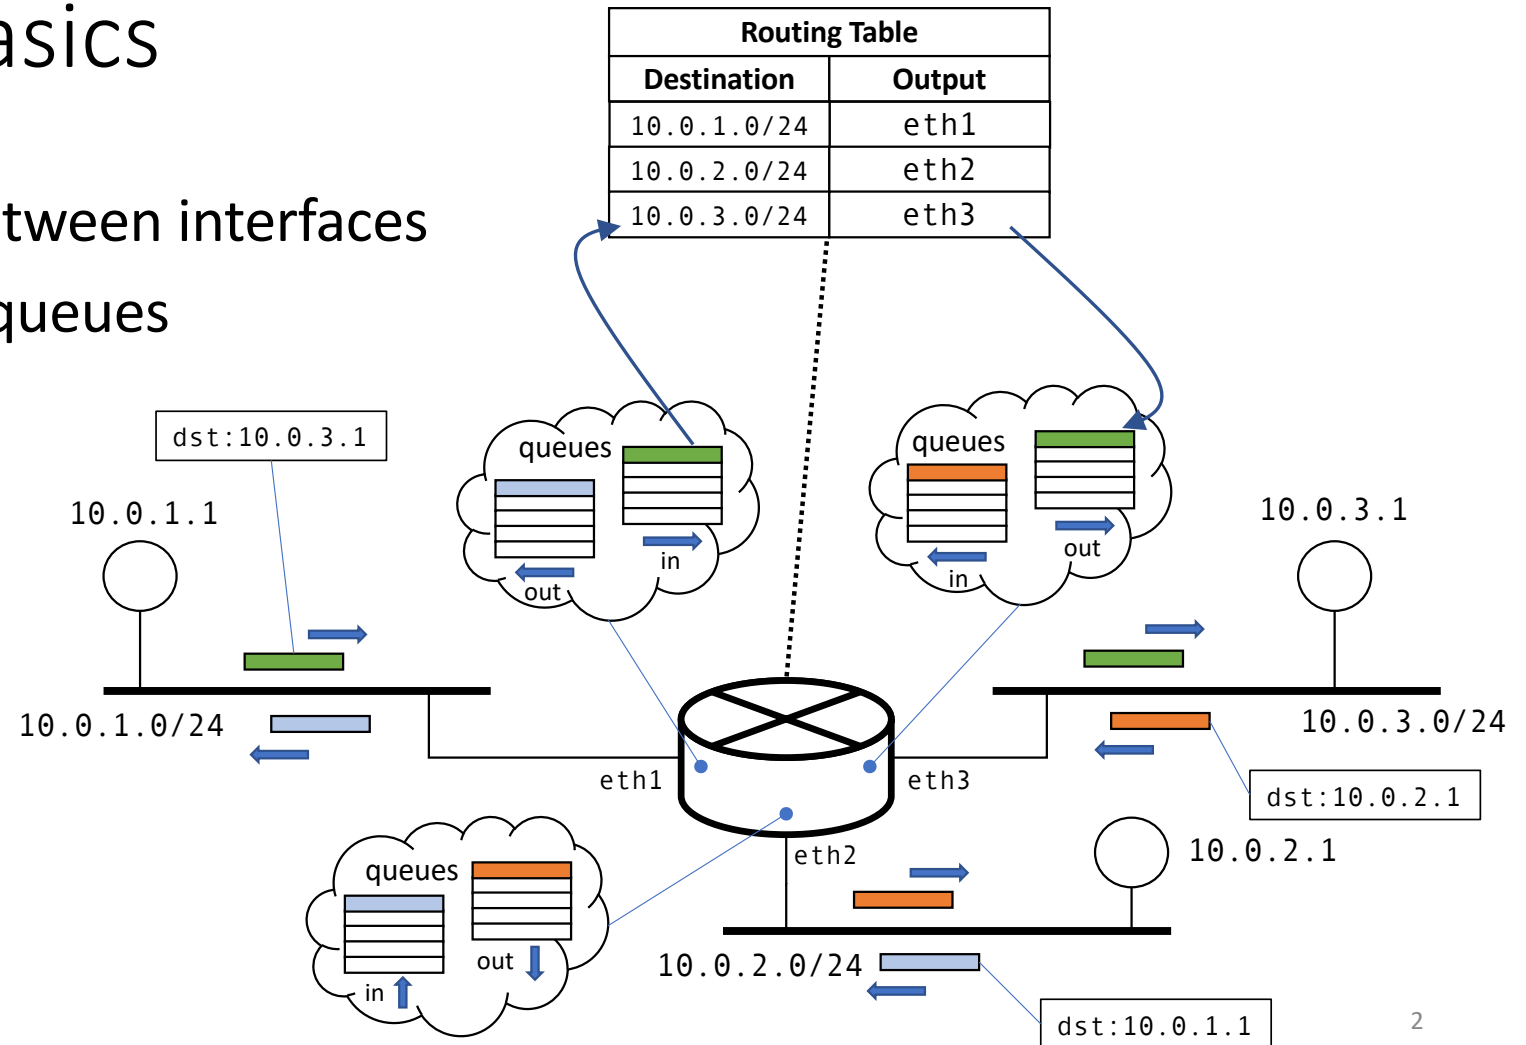
Routing

Ricardo Morla

FEUP – GORS/M.EEC, GRS/M.EIC

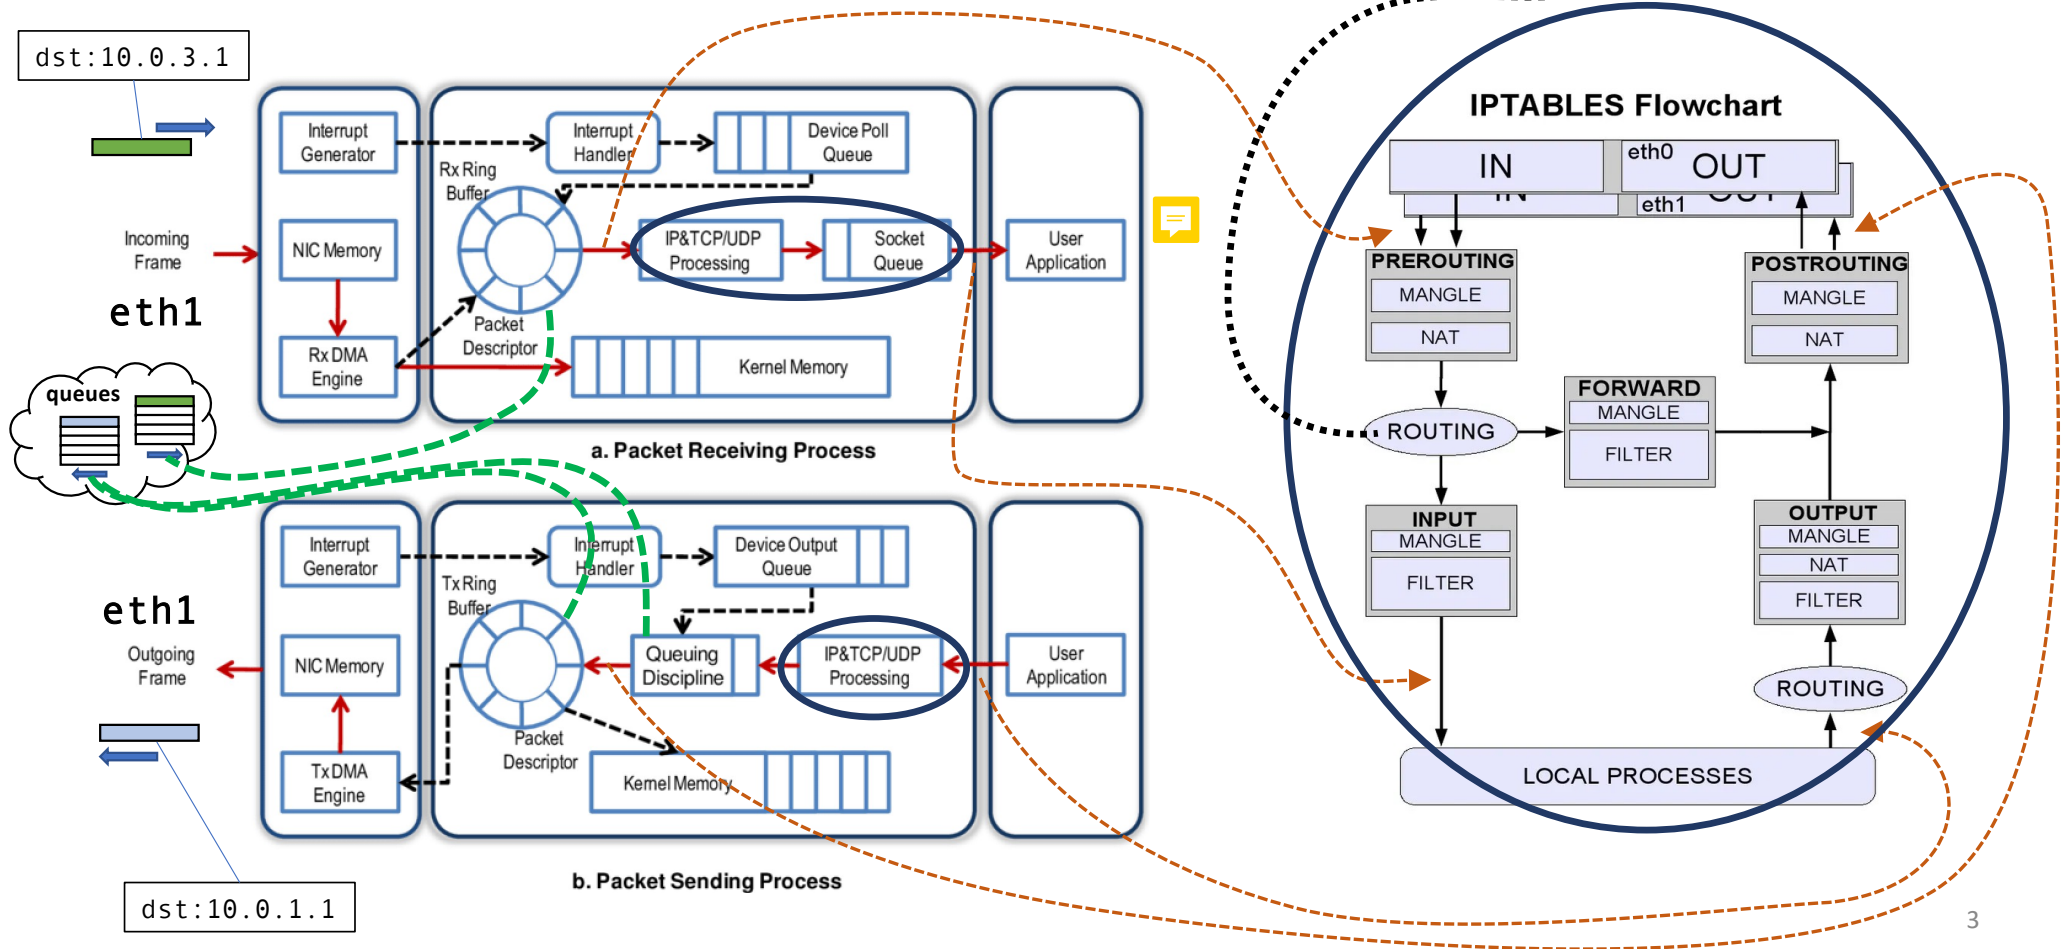
Routing basics

- Forwarding between interfaces
- Input/output queues
- Routing table



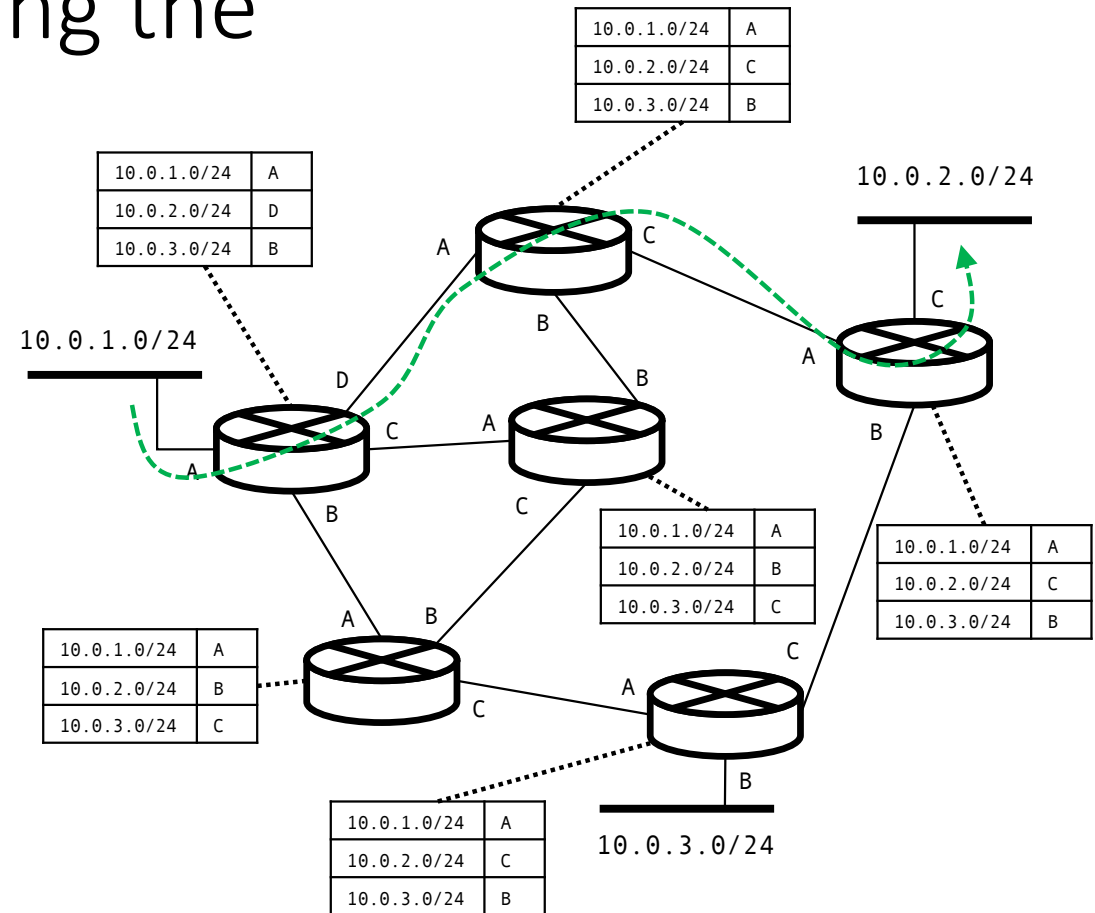
Inside the router

| Routing Table | |
|---------------|--------|
| Destination | Output |
| 10.0.1.0/24 | eth1 |
| 10.0.2.0/24 | eth2 |
| 10.0.3.0/24 | eth3 |




Routing – as in “choosing the paths for packets”

- Each router has its own routing table
- If routing tables are coherent, packets can go from any node in the network to any other node
- If they're not:
 - Loops
 - Blackholes




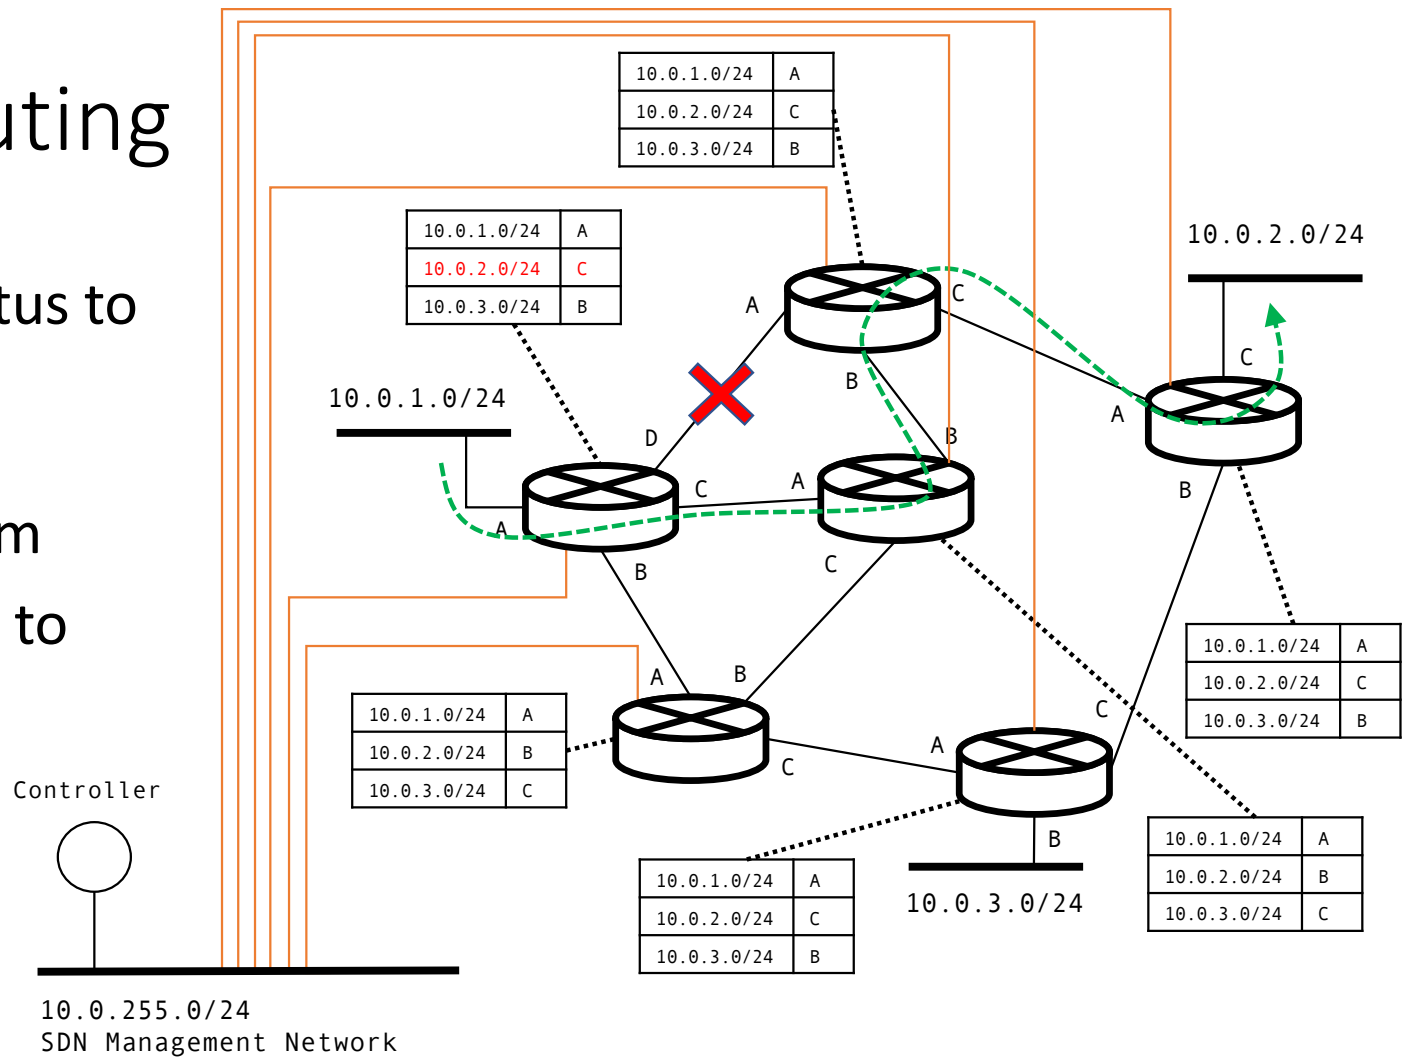
Reacting to changes in topology

- Each router can sense if their links are up or down
 - Or if new networks are added
- Automatically update routing tables
 - New path bypasses down link
- Based on some objective to optimize:
 - Smallest number of hops
 - Smallest delay
 - Largest throughput
 - Balance load on links
 - Administrative distance (...)




Centralized routing

- Nodes report link status to central location 
- Central location runs optimization algorithm
- Returns routing table to each router
- SDN-like solution
 - Data centers






Distributed routing

- Update routing tables dynamically
- According to information from other nodes
 - Each router knows the networks it is directly connected to
 - Routers send **control packets** to other routers
- Geographically dispersed network



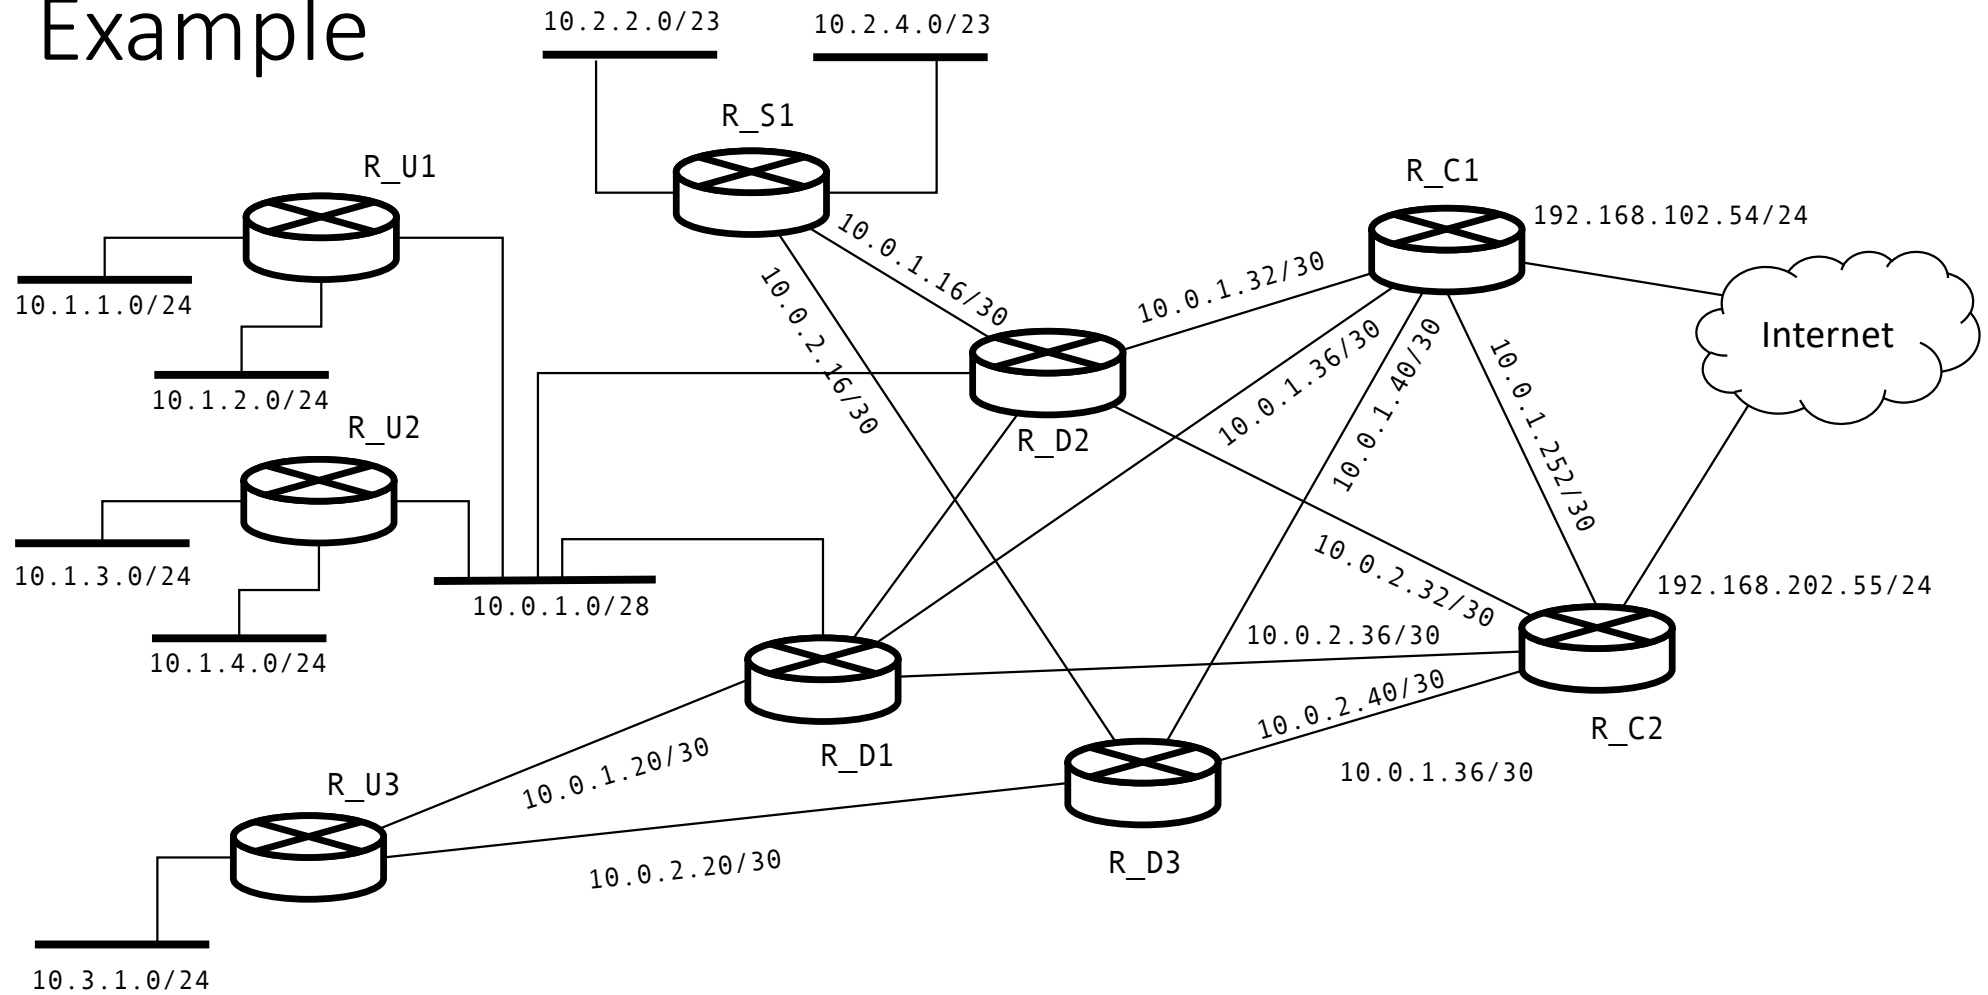
Routing algorithms

- Link state 
 - Full topology of the network, then Dijkstra
 - Distributed: flood link state announcements – OSPF 
- Distance vector
 - Distance to destination node, Bellman-Ford
 - Distributed: update distances, resend to neighbors
- Path vector 
 - List of nodes to destination node, local preferences
 - Distributed: update path vectors, resend to neighbors

Interior routing

- Same administrative domain
 - All networks and devices administered by same entity
 - Typically the internal network of an organization
- Protocols
 - RIP (outdated) – distance vector
 - EIGRP (Cisco) – distance vector, optimized, hybrid
 - OSPF – link state

Example

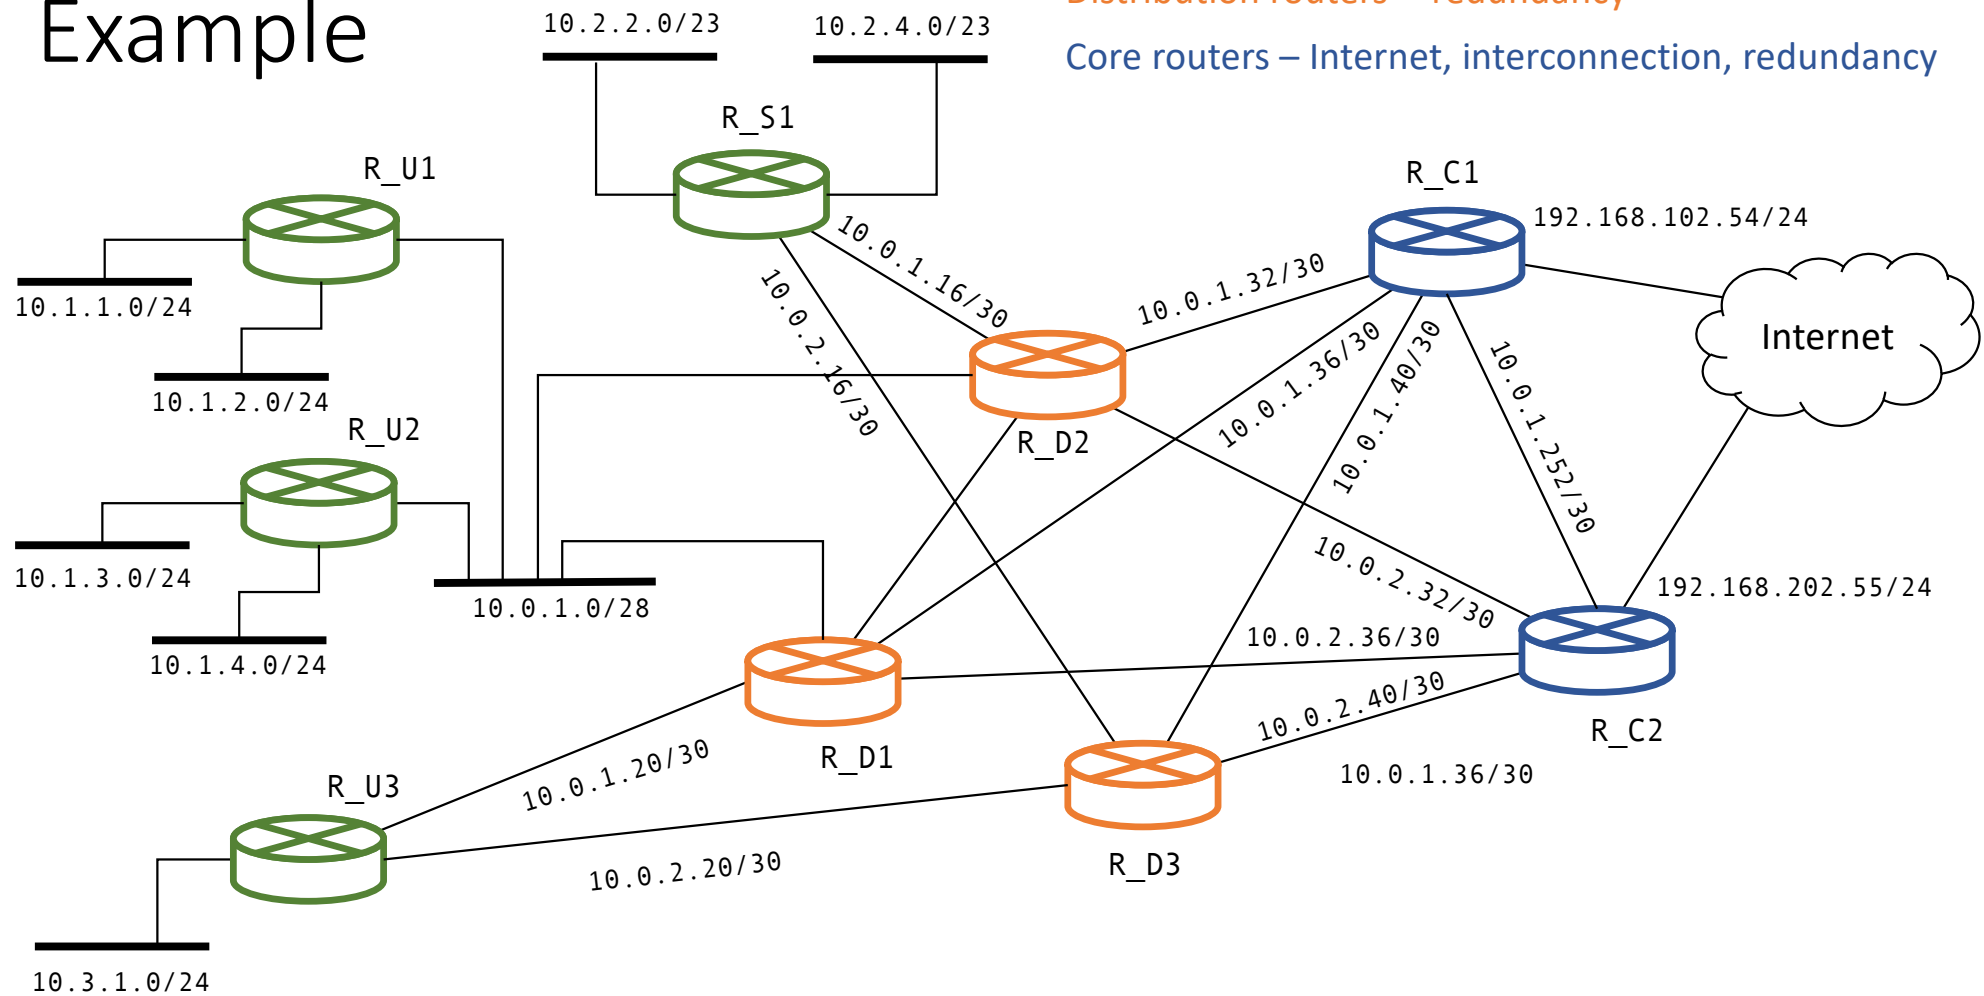


Example

Access routers – end user traffic to Internet and servers



Distribution routers -- redundancy

Core routers – Internet, interconnection, redundancy



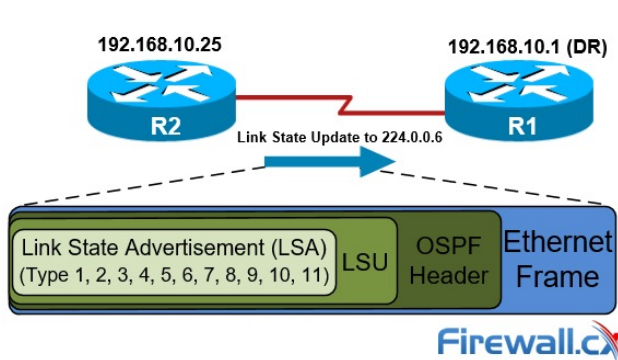


OSPF

- Messages

- Hello packets between routers to establish adjacency 
- Link state advertisement, link-state database 

- Designated routers, network adjacency vs. point-to-point link
- Cost metric manually defined, or multiple of reference bandwidth
- Backbone area, separate link-state database, Area Border Router

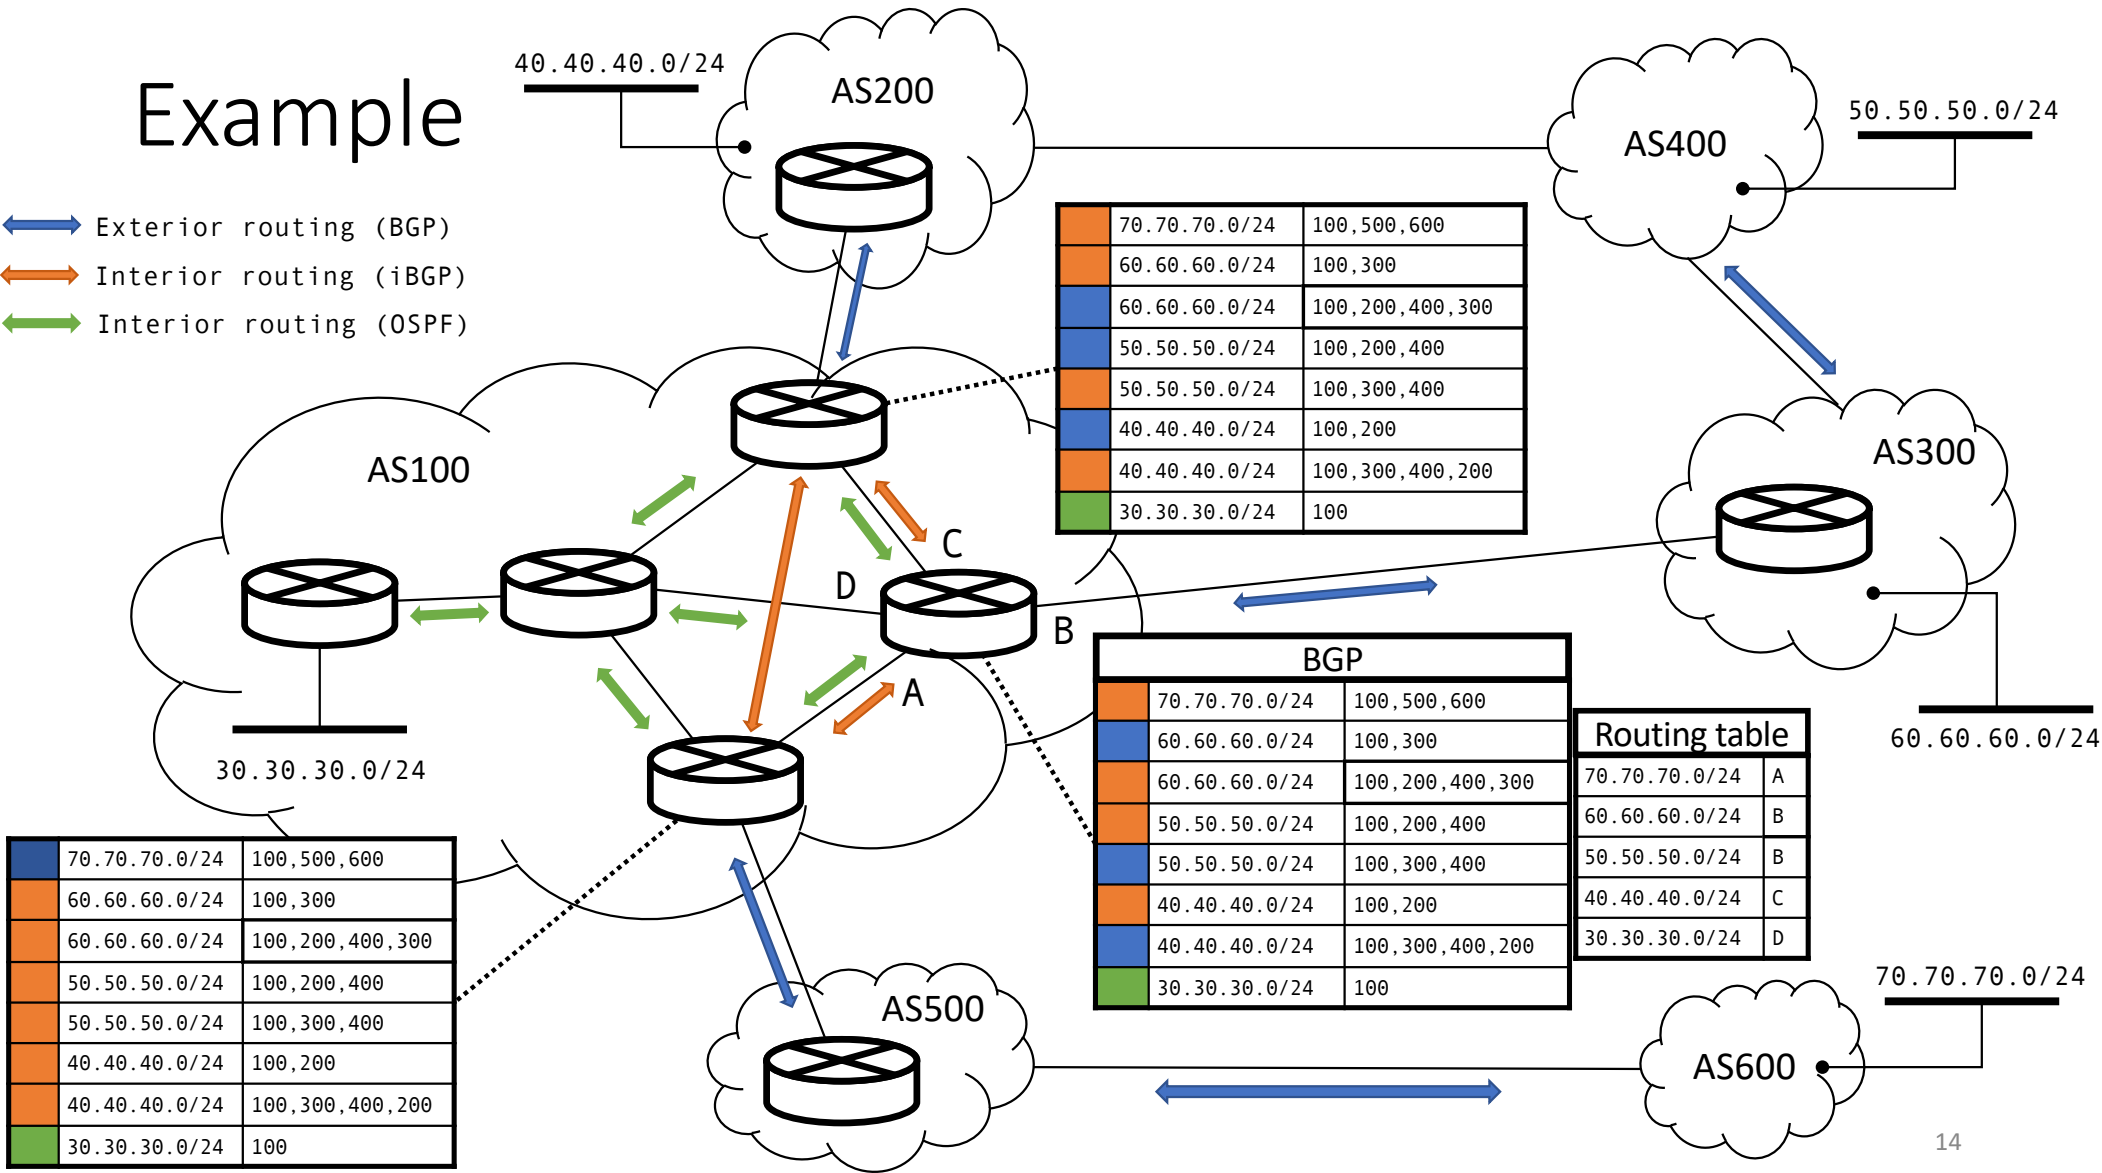


Exterior routing

- Different administrative domains
 - Each node in this network has a different administrator
 - The Internet
 - Each node is an Autonomous System – and a network by itself
- Protocols
 - BGP – path vector

Example




- ↔ Exterior routing (BGP)
- ↔ Interior routing (iBGP)
- ↔ Interior routing (OSPF)



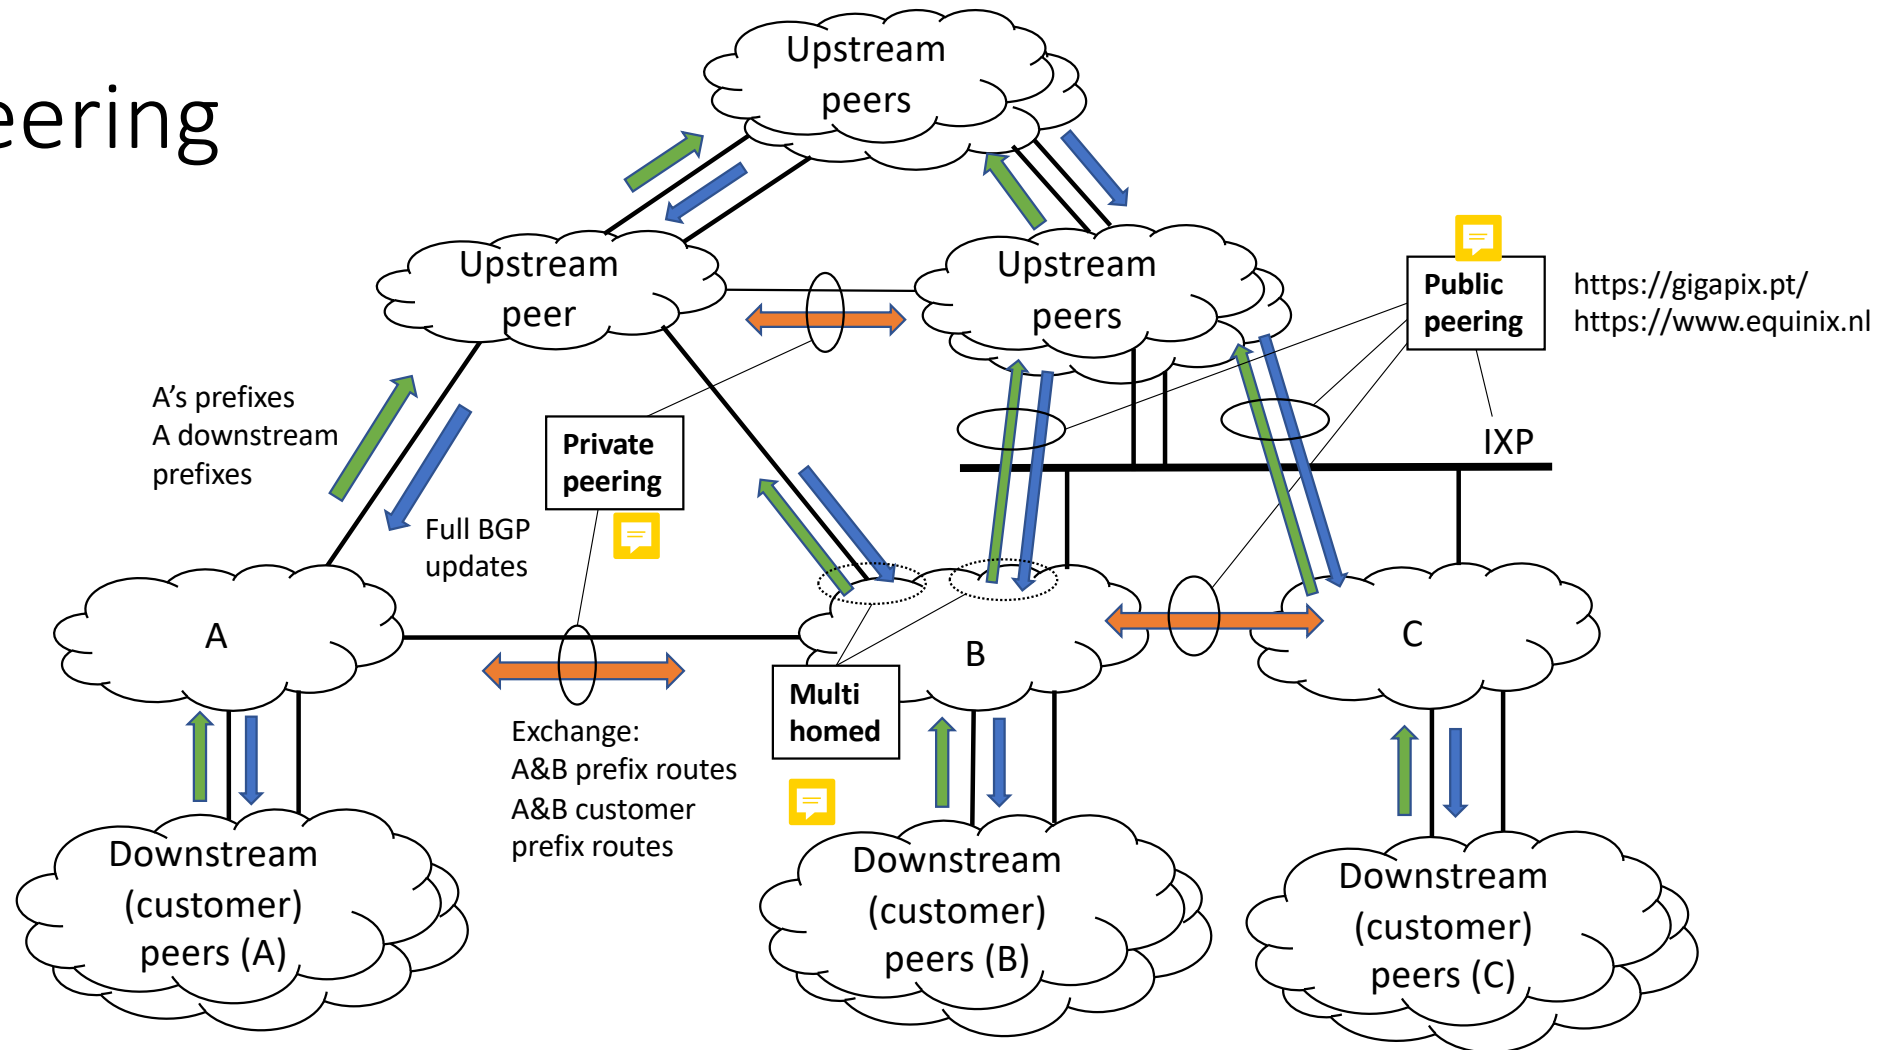
BGP protocol

- RFC 4271, BGP4
- Keep-alive messages, BGP update messages
- Exchange path-vector routes to other BGP routers
 - For each network aka network prefix
 - Select one route, announce it to neighbors, insert it in routing table
- Route map: set of rules to check what to do with route
 - Drop route, modify route and add it to routing table
- Route selection criteria
 - Weight (local to each router); AS-local preference for route (local to each AS), local route; length of AS path; MED multiple exit hint to other AS's, ...

BGP protocol

- TCP port 179
- Authentication, RFC 2385 
 - MD5 hash of password and BGP data
- IP header TTL set to 1 
 - Or more, if BGP connection over multiple IP hops
- Route aggregation 
 - 32.32.32.0/24, 32.32.33.0/24 announced as 32.32.32.0/23
- eBGP vs. iBGP
 - e: between routers of different AS
 - i: between routers of same AS in multiple peer AS's, full mesh, route reflectors

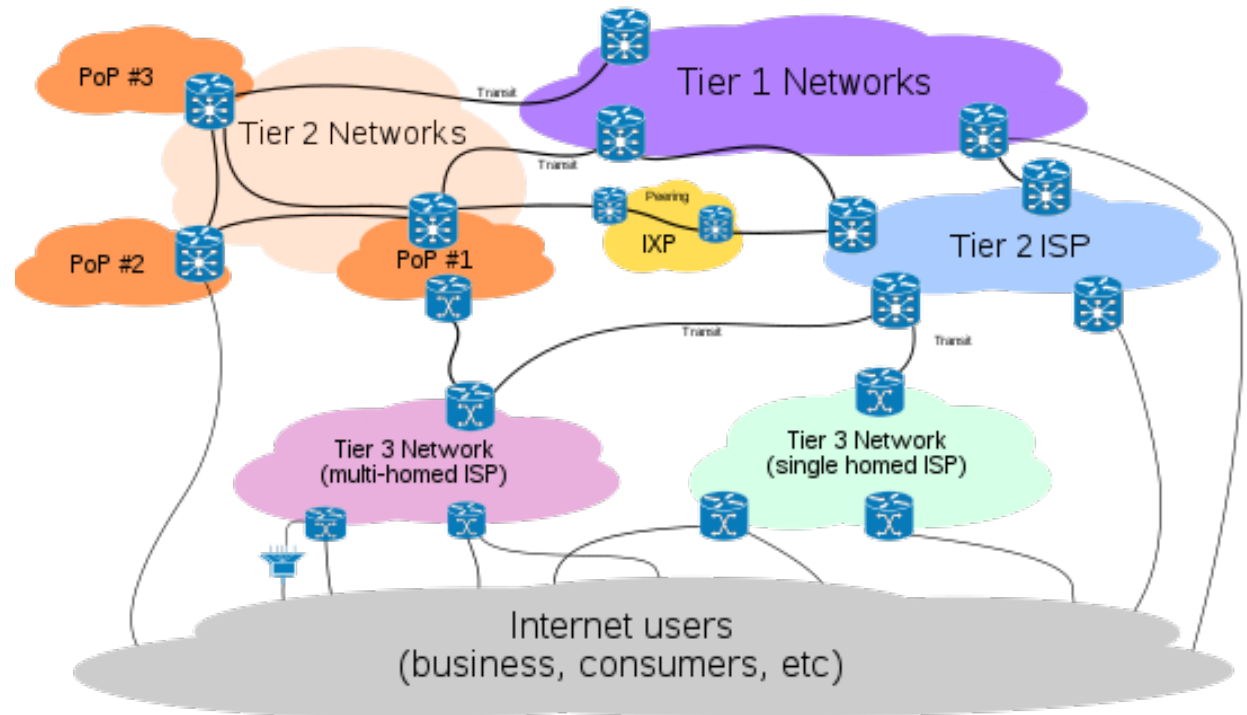
Peering



<http://www.bgphelp.com/2017/01/05/bgp-types-of-peering/>

Internet Tiers

- Tier 1
 - Settlement-free peering only
- Tier 2
 - Some free, other paying
- Tier 3
 - Paying only peering




Looking glass

- <http://www.routeviews.org>
- <https://bgp.he.net>
- RCCN
 - AS 1930
 - <https://bgp.he.net/AS1930>
- Archive
 - RIBS and UPDATES
- pyasn

```
import pyasn
asndb = pyasn.pyasn('ipasn_20140513.dat')
asndb.lookup('8.8.8.8')
```

https://bgp.he.net/AS1930#_graph4

AS1930 IPv4 Route Propagation



Management and Operations of Networks, Services, and Systems

Routing

Ricardo Morla

FEUP – GORS/M.EEC, GRS/M.EIC